

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT

on

Big Data Analytics (23CS6PCBDA)

Submitted by

Nithin Koushik P.V(1BM22CS284)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING

in

COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU-560019

Feb-2025 to June-2025

B. M. S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “**Big Data Analytics (23CS6PCBDA)**” carried out by **Nithin Koushik PV(1BM22CS185)**, who is bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the **Visvesvaraya Technological University, Belgaum** during the year 2024. The Lab report has been approved as it satisfies the academic requirements in respect of a **Big Data Analytics - (23CS6PCBDA)** work prescribed for the said degree.

Prof. Vikranth BM
Assistant Professor
BMSCE, Bengaluru

Dr. Kavitha Sooda
Professor and HoD of CSE
BMSCE, Bengaluru

Index Sheet

Sl. No.	Experiment Title	Page No.
1	MongoDB- CRUD Demonstration.	1
2	Perform the following DB operations using Cassandra. a) Create a keyspace by name Employee b) Create a column family by name Employee-Info with attributes Emp_Id Primary Key, Emp_Name, Designation, Date_of_Joining, Salary,Dept_Name c) Insert the values into the table in batch d) Update Employee name and Department of Emp-Id 121 e) Sort the details of Employee records based on salary f) Alter the schema of the table Employee_Info to add a column Projects which stores a set of Projects done by the corresponding Employee. g) Update the altered table to add project names. h) Create a TTL of 15 seconds to display the values of Employees.	5
3	Perform the following DB operations using Cassandra. a) Create a keyspace by name Library b) Create a column family by name Library-Info with attributes Stud_Id Primary Key, Counter_value of type Counter, Stud_Name, Book-Name, Book-Id, Date_of_issue c) Insert the values into the table in batch d) Display the details of the table created and increase the value of the counter e) Write a query to show that a student with id 112 has taken a book "BDA" 2 times. f) Export the created column to a csv file g) Import a given csv dataset from local file system into Cassandra column family	7
4	Execution of HDFS Commands for interaction with Hadoop Environment. (Minimum 10 commands to be executed)	9
5	Implement Wordcount program on Hadoop framework	11
6	From the following link extract the weather data https://github.com/tomwhite/hadoop-book/tree/master/input/ncdc/all Create a Map Reduce program to a) find average temperature for each year from NCDC data set. b) find the mean max temperature for every month.	15

7	For a given Text file, Create a Map Reduce program to sort the content in an alphabetic order listing only top 10 maximum occurrences of words.	24
8	Write a Scala program to print numbers from 1 to 100 using for loop.	30

9	Using RDD and FlatMap count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark.	31
---	--	----

3

Course Outcome

CO 1	Apply the concept of NoSQL, Hadoop or Spark for a given task
CO 2	Analyze big data analytics mechanisms that can be applied to obtain solution for a given problem.
CO 3	Design and implement solutions using data analytics mechanisms for a given problem.

4

Experiment-1

Q) MongoDB- CRUD Operations Demonstration (Practice and Self Study)

Code & Output:

1. Create a database “Student” with the following attributes Rollno, Name , Age, ContactNo, Email-Id, grade, hobby:

use Students;

2. Insert 5 appropriate values according to the below queries.

```
db.students.insertMany([  
  { "Rollno": 10, "Name": "John", "Age": 20, "ContactNo": "1234567890",  
    "Email-Id": "john@example.com", "grade": "A", "hobby": "Reading" },  
  { "Rollno": 11, "Name": "Alice", "Age": 21, "ContactNo": "9876543210",  
    "Email-Id": "alice@example.com", "grade": "B", "hobby": "Painting" },  
  { "Rollno": 12, "Name": "Bob", "Age": 22, "ContactNo": "2345678901", "Email-Id":  
    "bob@example.com", "grade": "C", "hobby": "Cooking" },  
  { "Rollno": 13, "Name": "Eve", "Age": 23, "ContactNo": "3456789012", "Email-Id":  
    "eve@example.com", "grade": "A" },  
  { "Rollno": 14, "Name": "Charlie", "Age": 24, "ContactNo": "4567890123",  
    "Email-Id": "charlie@example.com", "hobby": "Gardening" }  
])
```

```

Atlas atlas-wanmtx-shard-0 [primary] Students> use Students
switched to db Students
Atlas atlas-wanmtx-shard-0 [primary] Students> show collections

Atlas atlas-wanmtx-shard-0 [primary] Students> db.students.insertMany([
...   { "Rollno": 10, "Name": "John", "Age": 20, "ContactNo": "1234567890", "Email-Id":
"john@example.com", "grade": "A", "hobby": "Reading" },
...   { "Rollno": 11, "Name": "Alice", "Age": 21, "ContactNo": "9876543210", "Email-Id":
"alice@example.com", "grade":
"B", "hobby": "Painting" },
...   { "Rollno": 12, "Name": "Bob", "Age": 22, "ContactNo": "2345678901", "Email-Id": "
bob@example.com", "grade": "C", "hobby": "Cooking" },
...   { "Rollno": 13, "Name": "Eve", "Age": 23, "ContactNo": "3456789012", "Email-Id": "
eve@example.com", "grade": "A"
},
...   { "Rollno": 14, "Name": "Charlie", "Age": 24, "ContactNo": "4567890123", "Email-Id
": "charlie@example.com", "hobby": "Gardening" }
... ])
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId("661ce9dc76a00ff8cc51dae1"),
    '1': ObjectId("661ce9dc76a00ff8cc51dae2"),
    '2': ObjectId("661ce9dc76a00ff8cc51dae3"),
    '3': ObjectId("661ce9dc76a00ff8cc51dae4"),
    '4': ObjectId("661ce9dc76a00ff8cc51dae5")
  }
}

```

3. Write query to update Email-Id of a student with rollno 10.

```

db.students.updateOne(
  { "Rollno": 10 },
  { $set: { "Email-Id": "john.doe@example.com" } }
)

```

```

Atlas atlas-wanmtx-shard-0 [primary] Students> db.students.updateOne(
...   { "Rollno": 10 },
...   { $set: { "Email-Id": "john.doe@example.com" } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}

```

4. Replace the student name from “Alice” to “Alicee” of rollno 11

```

db.students.updateOne(
  { "Rollno": 11 },
  { $set: { "Name": "Alicee" } }
)

```

```
Atlas atlas-wanmtx-shard-0 [primary] Students> db.students.updateOne(
...   { "Rollno": 11 },
...   { $set: { "Name": "Alice" } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

5. Display Student Name and grade(Add if grade is not present)where the _id column is 1.

```
db.students.find({}, { "Name": 1, "grade": { $ifNull: ["$grade", "Not available"] }, "_id": 0 })
```

```
Atlas atlas-wanmtx-shard-0 [primary] Students> db.students.find({}, { "Name": 1, "grade":
{ $ifNull: ["$grade", "Not available"] }, "_id": 0 })
[
  { Name: 'John', grade: 'A' },
  { Name: 'Alice', grade: 'B' },
  { Name: 'Bob', grade: 'C' },
  { Name: 'Eve', grade: 'A' },
  { Name: 'Charlie', grade: 'Not available' }
]
```

6. Update to add hobbies

```
db.students.updateMany(
  { "Name": "Eve" },
  { $set: { "hobby": "Dancing" } }
)
```

```
Atlas atlas-wanmtx-shard-0 [primary] Students> db.students.updateMany(
...   { "Name": "Eve" },
...   { $set: { "hobby": "Dancing" } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

7. Find documents where hobbies is set neither to Chess nor to Skating

```
db.students.find({ "hobby": { $nin: ["Chess", "Skating"] } })
```

```
Atlas atlas-wanmtx-shard-0 [primary] Students> db.students.find({ "hobby": { $nin: ["Chess", "Skating"] } })
[
  {
    _id: ObjectId("661ce9dc76a00ff8cc51dae1"),
    Rollno: 10,
    Name: 'John',
    Age: 20,
    ContactNo: '1234567890',
    'Email-Id': 'john.doe@example.com',
    grade: 'A',
    hobby: 'Reading'
  },
  {
    _id: ObjectId("661ce9dc76a00ff8cc51dae2"),
    Rollno: 11,
    Name: 'Alice',
    Age: 21,
    ContactNo: '9876543210',
    'Email-Id': 'alice@example.com',
    grade: 'B',
    hobby: 'Painting'
  },
  {
    _id: ObjectId("661ce9dc76a00ff8cc51dae3"),
    Rollno: 12,
    Name: 'Bob',
    Age: 22,
    ContactNo: '2345678901',
    'Email-Id': 'bob@example.com',
    grade: 'C',
    hobby: 'Cooking'
  }
]
```

8. Find documents whose name begins with A

```
db.students.find({ "Name": /^A/ })
```

```
Atlas atlas-wanmtx-shard-0 [primary] Students> db.students.find({ "Name": /^A/ })
[
  {
    _id: ObjectId("661ce9dc76a00ff8cc51dae2"),
    Rollno: 11,
    Name: 'Alice',
    Age: 21,
    ContactNo: '9876543210',
    'Email-Id': 'alice@example.com',
    grade: 'B',
    hobby: 'Painting'
  }
]
```


Experiment – 2

Q) Perform the following DB operations using Cassandra

- Create a keyspace by name **Employee**
- Create a column family by name **Employee-Info** with attributes
Emp_Id Primary Key, Emp_Name,
Designation, Date_of_Joining, Salary, Dept_Name
- Insert the values into the table in **batch**
- Update Employee name and Department of **Emp-Id 121**
- Sort the details of Employee records based on **salary**
- Alter the schema of the table **Employee_Info** to add a column **Projects**
which stores a **set of Projects** done by the corresponding Employee.
- Update the altered table to **add project names**
- Create a **TTL of 15 seconds** to display the values of Employees

Code & Output:

```
bruce@bruce-HP-Elite-Tower-860-G8-Desktop-PC: $ cqlsh
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.1.4 | CQL spec 3.4.6 | Native protocol v5]
Use HELP for help.
cqlsh> create keyspace Employee with replication = {'class': 'SimpleStrategy', 'replication_factor': 1};
SyntaxException: line 1:80 mismatched input ';' expecting ')' (...with replication = {'class': 'SimpleStrategy', 'replication_factor': 1};...)
cqlsh> create keyspace Employee WITH replication={'class': 'SimpleStrategy', 'replication_factor': 1};
ConfigurationException: unrecognized strategy option (replication_factor) passed to SimpleStrategy for keyspace employee
cqlsh> create keyspace Employee WITH replication={'class': 'SimpleStrategy', 'replication_factor': 1};
cqlsh> DESCRIBE KEYSPACES

employee system_auth          system_schema system_views
system    system_distributed system_traces system_virtual_schema

cqlsh> CREATE TABLE IF NOT EXISTS Employee_Info(
... Emp_Id INT PRIMARY KEY,
... Emp_name TEXT,
... designation TEXT,
... date_of_joining DATE,
... Salary FLOAT,
... Dep_name TEXT,
... Projects SET<TEXT>);
InvalidRequest: Error from server: code=2200 [Invalid query] message="No keyspace has been specified. USE a keyspace, or explicitly specify keyspace.tablename"
cqlsh> USE EMPLOYEE
...
cqlsh> USE Employee
...
cqlsh> USE Employee;
cqlsh:employee> CREATE TABLE IF NOT EXISTS Employee_Info( Emp_Id INT PRIMARY KEY, Emp_name TEXT, designation TEXT, date_of_joining DATE, Salary FLOAT, Dep_name TEXT, Projects SET<TEXT>);
cqlsh:employee> describe keyspace Employee

CREATE KEYSPACE employee WITH replication = {'class': 'SimpleStrategy', 'replication_factor': '1'} AND durable_writes = true;

CREATE TABLE employee.employee_info (
  emp_id int PRIMARY KEY,
  date_of_joining date,
  dep_name text,
  designation text,
  emp_name text,
  salary float,
  projects set<text>
) WITH additional_write_policy = '99p'
AND bloom_filter_fp_chance = 0.01
AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
AND cdc = false
AND comment = ''
AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
AND crc_check_chance = 1.0
AND default_time_to_live = 0
AND extensions = {}
AND gc_grace_seconds = 864000
AND max_index_interval = 2549
AND memtable_flush_period_in_ms = 0
AND min_index_interval = 128
```

```

cqlsh:employee> update employee_info using ttl 15 set salary = 0 where emp_id = 121;
cqlsh:employee> select * from employee_info;

```

emp_id	bonus	date_of_joining	dep_name	designation	emp_name	projects	salary
120	12000	2024-05-06	Engineering	Developer	Priyanka GH	['Project B', 'ProjectA']	1e+06
123	null	2024-05-07	Engineering	Engineer	Sadhana	['Project M', 'Project P']	1.2e+06
122	null	2024-05-06	Management	HR	Rachana	['Project C', 'Project M']	9e+05
121	11000	2024-05-06	Management	Developer	Shreya	['Project C', 'ProjectA']	0

(4 rows)

```
cqlsh:employee> select * from employee_info;
```

emp_id	bonus	date_of_joining	dep_name	designation	emp_name	projects	salary
120	12000	2024-05-06	Engineering	Developer	Priyanka GH	['Project B', 'ProjectA']	1e+06
123	null	2024-05-07	Engineering	Engineer	Sadhana	['Project M', 'Project P']	1.2e+06
122	null	2024-05-06	Management	HR	Rachana	['Project C', 'Project M']	9e+05
121	11000	2024-05-06	Management	Developer	Shreya	['Project C', 'ProjectA']	null

(4 rows)

```
cqlsh:employee>
```

```
AND speculative_retry = '99p';
```

```
cqlsh:employee> select * from employee_info;
```

emp_id	date_of_joining	dep_name	designation	emp_name	projects	salary
120	2024-05-06	Engineering	Developer	Priyanka	['Project B', 'ProjectA']	1e+06
123	2024-05-07	Engineering	Engineer	Sadhana	['Project M', 'Project P']	1.2e+06
122	2024-05-06	Management	HR	Rachana	['Project C', 'Project M']	9e+05
121	2024-05-06	Management	Developer	Shreya	['Project C', 'ProjectA']	9e+05

(4 rows)

```
cqlsh:employee> update employee_info set emp_name = 'Priyanka GH' where emp_id = '120';
```

```
InvalidRequest: Error from server: code=2200 [Invalid query] message="Invalid STRING constant (120) for "emp_id" of type int"
```

```
cqlsh:employee> update employee_info set emp_name = 'Priyanka GH' where emp_id=120;
```

```
cqlsh:employee> select * from employee_info;
```

emp_id	date_of_joining	dep_name	designation	emp_name	projects	salary
120	2024-05-06	Engineering	Developer	Priyanka GH	['Project B', 'ProjectA']	1e+06
123	2024-05-07	Engineering	Engineer	Sadhana	['Project M', 'Project P']	1.2e+06
122	2024-05-06	Management	HR	Rachana	['Project C', 'Project M']	9e+05
121	2024-05-06	Management	Developer	Shreya	['Project C', 'ProjectA']	9e+05

(4 rows)

```
cqlsh:employee> select * from employee_info order by salary;
```

```
InvalidRequest: Error from server: code=2200 [Invalid query] message="ORDER BY is only supported when the partition key is restricted by an EQ or an IN."
```

```
cqlsh:employee> alter table employee_info add bonus INT;
```

```
cqlsh:employee> select * from employee_info;
```

emp_id	bonus	date_of_joining	dep_name	designation	emp_name	projects	salary
120	null	2024-05-06	Engineering	Developer	Priyanka GH	['Project B', 'ProjectA']	1e+06
123	null	2024-05-07	Engineering	Engineer	Sadhana	['Project M', 'Project P']	1.2e+06
122	null	2024-05-06	Management	HR	Rachana	['Project C', 'Project M']	9e+05
121	null	2024-05-06	Management	Developer	Shreya	['Project C', 'ProjectA']	9e+05

(4 rows)

```
cqlsh:employee> update employee_info set bonus = 12000 where emp_id = 120;
```

```
cqlsh:employee> select * from employee_info;
```

emp_id	bonus	date_of_joining	dep_name	designation	emp_name	projects	salary
120	12000	2024-05-06	Engineering	Developer	Priyanka GH	['Project B', 'ProjectA']	1e+06
123	null	2024-05-07	Engineering	Engineer	Sadhana	['Project M', 'Project P']	1.2e+06
122	null	2024-05-06	Management	HR	Rachana	['Project C', 'Project M']	9e+05
121	null	2024-05-06	Management	Developer	Shreya	['Project C', 'ProjectA']	9e+05

(4 rows)

```
cqlsh:employee> update employee_info set bonus = 11000 where emp_id = 121;
```

```
cqlsh:employee> select * from employee_info using ttl 15 where emp_id = 123;
```

```
SyntaxException: line 1:28 mismatched input 'using' expecting EOF (select * from employee_info [using] ttl...)
```

```
cqlsh:employee> select * from employee_info where emp_id = 121 using ttl 15;
```

```
SyntaxException: line 1:47 no viable alternative at input 'using' (...employee_info where emp_id = 121 [using]...)
```

```
cqlsh:employee> update employee_info using ttl 15 set salary = 0 where emp_id = 121;
```

```
cqlsh:employee> select * from employee_info;
```

Experiment – 3

Q) Perform the following DB operations using Cassandra

- Create a keyspace by name **Library**
- Create a column family by name **Library-Info** with attributes
Stud_Id Primary Key,
Counter_value of type **Counter,**
Stud_Name, Book-Name, Book-Id,
Date_of_issue
- Insert the values into the table in **batch**
- Display the details of the table created and **increase the value of the counter**
- Write a query to show that a student with **id 112** has taken a book “**BDA**” **2 times** f) **Export** the created column to a **CSV file**
- Import** a given CSV dataset from **local file system** into Cassandra **column family** Code & Output:

```
mscscse@mscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ cqlsh
Connected to Test Cluster at 127.0.0.1:9042
(cqlsh 6.1.0 | Cassandra 4.1.4 | CQL spec 3.4.0 | Native protocol v5)
Use HELP for help.
cqlsh> CREATE KEYSPACE students WITH REPLICATION={
... 'class':'SimpleStrategy','replication_factor':1};
cqlsh> DESCRIBE KEYSPACES

students  system_auth          system_schema  system_views
system    system_distributed        system_traces  system_virtual_schema

cqlsh> SELECT * FROM system.schema_keyspaces;
InvalidRequest: Error from server: code=2200 [Invalid query] message="table schema_keyspaces does not exist"
cqlsh> use students;
cqlsh:students> create table Students_info(Roll_No int Primary key,StudName text,DateOfJoining timestamp,last_exam_Percent double);
cqlsh:students> describe tables;

students_info

cqlsh:students> describe table students;
Table 'students' not found in keyspace 'students'
cqlsh:students> describe table students_info;

CREATE TABLE students.students_info (
  roll_no int PRIMARY KEY,
  dateofjoining timestamp,
  last_exam_percent double,
  studname text
) WITH additional_write_policy = '99p'
  AND bloom_filter_fp_chance = 0.01
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
  AND cdc = false
  AND comment = ''
  AND compaction = [{'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}]
  AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
  AND crc_check_chance = 1.0
  AND default_time_to_live = 0
  AND extensions = {}
  AND gc_grace_seconds = 864000
  AND max_index_interval = 2048
  AND memtable_flush_period_in_ms = 0
  AND min_index_interval = 128
  AND read_repair = 'BLOCKING'
  AND speculative_retry = '99p';
```

```

[4 rows]
sqlsh:students> select * from students_info where roll_no in (1,2,3);

roll_no | dateofjoining          | last_exam_percent | studname
-----|-----|-----|-----
1 | 2021-10-08 18:30:00.000000+0000 | 98 | Sidhans
2 | 2021-10-09 18:30:00.000000+0000 | 97 | Rutu
3 | 2021-10-09 18:30:00.000000+0000 | 97.5 | Rashans

[3 rows]
sqlsh:students> select * from students_info where Studname='Charu';
InvalidRequest: Error from server: code=2280 [Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query it requires the performance_sensitivity hint. See 81106 P31708390"
sqlsh:students> create index on Students_info(Studname);
sqlsh:students> select * from students_info where Studname='Charu';

roll_no | dateofjoining          | last_exam_percent | studname
-----|-----|-----|-----
4 | 2021-10-05 18:30:00.000000+0000 | 96.5 | Charu

[1 rows]
sqlsh:students> select Roll_no,Studname from students_info LIMIT 2;

roll_no | studname
-----|-----
1 | Sidhans
2 | Rutu

[2 rows]
sqlsh:students> SELECT Roll_no as 'IDN' from Students_info;

IDN
---
1
2
4
3

```


Experiment - 4

Q) Execution of HDFS Commands for interaction with Hadoop Environment. (Minimum 10 commands to be executed)

Code & Output:

```
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ cd ./Desktop/
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [bmscecse-HP-Elite-Tower-800-G9-Desktop-PC]
Starting resourcemanager
Starting nodemanagers
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hdfs dfs -mkdir /Lab05
```

```
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ touch test.txt
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ nano text.txt
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hdfs dfs -put ./text.txt /Lab05/text.txt
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hadoop fs -ls /Lab05
Found 1 items
-rw-r--r-- 1 hadoop supergroup      19 2024-05-13 14:33 /Lab05/text.txt
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hdfs dfs -cat /Lab05/text.txt
Hello
How are you?
```

```
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hadoop fs -ls /Lab05
Found 2 items
-rw-r--r-- 1 hadoop supergroup      15 2024-05-13 14:40 /Lab05/test.txt
-rw-r--r-- 1 hadoop supergroup      19 2024-05-13 14:33 /Lab05/text.txt
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hdfs dfs -getmerge /Lab05 /text.txt /Lab05 /test.txt ../Downloads/Merged.txt
getmerge: `/text.txt': No such file or directory
getmerge: `/test.txt': No such file or directory
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hdfs dfs -getmerge /Lab05/text.txt /Lab05/test.txt ../Downloads/Merged.txt
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hadoop fs -getfacl /Lab05
# file: /Lab05
# owner: hadoop
# group: supergroup
user::rwx
group::r-x
other::r-x
```

Experiment - 5

Q) Implement Wordcount program on Hadoop framework

Code & Output:

Mapper Code: WCMapper.java

java

CopyEdit

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

public class WCMapper extends MapReduceBase implements Mapper<LongWritable, Text, Text,
IntWritable> {

    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter rep)
throws IOException {

        String line = value.toString();
        for (String word : line.split(" ")) {
            if (word.length() > 0) {
                output.collect(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

Reducer Code: WCReducer.java

java

CopyEdit

```

import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class WCReducer extends MapReduceBase implements Reducer<Text, IntWritable, Text,
IntWritable> {

    public void reduce(Text key, Iterator<IntWritable> value, OutputCollector<Text, IntWritable> output,
Reporter rep) throws IOException {

        int count = 0;

        while (value.hasNext()) {
            IntWritable i = value.next();
            count += i.get();
        }
        output.collect(key, new IntWritable(count));
    }
}

```

Driver Code: WCDriver.java

java

CopyEdit

```

import java.io.IOException;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapred.FileInputFormat;

```

```

import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class WCDriver extends Configured implements Tool {
    public int run(String args[]) throws IOException {
        if (args.length < 2) {
            System.out.println("Please give valid inputs");
            return -1;
        }

        JobConf conf = new JobConf(WCDriver.class);
        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));

        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class);
        conf.setMapOutputKeyClass(Text.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);

        JobClient.runJob(conf);
        return 0;
    }

    public static void main(String args[]) throws Exception {
        int exitCode = ToolRunner.run(new WCDriver(), args);

        System.out.println(exitCode);
    }
}

```


}

}

Input File -> big data hadoop big data analytics

map reduce big data

Output:

(big, 1)

(data, 1)

(hadoop, 1)

(big, 1)

(data, 1)

(analytics, 1)

(map, 1)

(reduce, 1)

(big, 1)

(data, 1)

Experiment – 6

Q) From the following link extract the weather data

<https://github.com/tomwhite/hadoopbook/tree/master/input/ncdc/all>

Create a Map Reduce program to

- a) find average temperature for each year from NCDC data set.
- b) find the mean max temperature for every month.

Code & Output:

- a) Find average temperature for each year from NCDC data set

AverageDriver.java

java

CopyEdit

```
package temp;
```

```
import org.apache.hadoop.fs.Path;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Job;
```

```
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
```

```
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
public class AverageDriver {
```

```
    public static void main(String[] args) throws Exception {
```

```
        if (args.length != 2) {
```

```
            System.err.println("Please Enter the input and output parameters");
```

```
            System.exit(-1);
```

```
        }
```

```
        Job job = new Job();
```

```
        job.setJarByClass(AverageDriver.class);
```

```
        job.setJobName("Max temperature");
```

```

FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));

job.setMapperClass(AverageMapper.class);
job.setReducerClass(AverageReducer.class);

job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

AverageMapper.java

java

CopyEdit

package temp;

```

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

```

```

public class AverageMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
public static final int MISSING = 9999;

```

```

public void map(LongWritable key, Text value, Mapper<LongWritable, Text, Text, IntWritable>.Context
context)

throws IOException, InterruptedException {

```

```

int temperature;

String line = value.toString();
String year = line.substring(15, 19);

if (line.charAt(87) == '+') {
    temperature = Integer.parseInt(line.substring(88, 92));
} else {
    temperature = Integer.parseInt(line.substring(87, 92));
}

String quality = line.substring(92, 93);
if (temperature != 9999 && quality.matches("[01459]"))
    context.write(new Text(year), new IntWritable(temperature));
}
}

```

AverageReducer.java

java

CopyEdit

package temp;

```

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class AverageReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values,

        Reducer<Text, IntWritable, Text, IntWritable>.Context context)
        throws IOException, InterruptedException {

```

```
int max_temp = 0;
int count = 0;

for (IntWritable value : values) {
    max_temp += value.get();
    count++;
}

context.write(key, new IntWritable(max_temp / count));
}
}
```

b) Find the mean max temperature for every month

MeanMaxDriver.java

java

CopyEdit

```
package meanmax;
```

```
import org.apache.hadoop.fs.Path;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Job;
```

```
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
```

```
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
public class MeanMaxDriver {
```

```

public static void main(String[] args) throws Exception {
    if (args.length != 2) {
        System.err.println("Please Enter the input and output parameters");
        System.exit(-1);
    }
}

```

23

```

Job job = new Job();
job.setJarByClass(MeanMaxDriver.class);
job.setJobName("Max temperature");

FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));

job.setMapperClass(MeanMaxMapper.class);
job.setReducerClass(MeanMaxReducer.class);

job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

MeanMaxMapper.java

```

java
CopyEdit
package meanmax;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;

```

```

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class MeanMaxMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
    public static final int MISSING = 9999;

```

24

```

    public void map(LongWritable key, Text value,
        Mapper<LongWritable, Text, Text, IntWritable>.Context context)
        throws IOException, InterruptedException {

```

```

        int temperature;
        String line = value.toString();
        String month = line.substring(19, 21);

```

```

        if (line.charAt(87) == '+') {
            temperature = Integer.parseInt(line.substring(88, 92));
        } else {
            temperature = Integer.parseInt(line.substring(87, 92));
        }

```

```

        String quality = line.substring(92, 93);
        if (temperature != 9999 && quality.matches("[01459]"))
            context.write(new Text(month), new IntWritable(temperature));
        }
    }
}

```

MeanMaxReducer.java

```
java
CopyEdit
package meanmax;

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class MeanMaxReducer extends Reducer<Text, IntWritable, Text, IntWritable> { 25
    public void reduce(Text key, Iterable<IntWritable> values,
        Reducer<Text, IntWritable, Text, IntWritable>.Context context)
        throws IOException, InterruptedException {

        int max_temp = 0;
        int total_temp = 0;
        int count = 0;
        int days = 0;

        for (IntWritable value : values) {
            int temp = value.get();
            if (temp > max_temp)
                max_temp = temp;

            count++;
            if (count == 3) {
                total_temp += max_temp;
                max_temp = 0;
                count = 0;
            }
        }
    }
}
```



```

days++;
}
}

context.write(key, new IntWritable(total_temp / days));
}
}

```

Map-Reduce Framework

```

Map input records=6
Map output records=6
Map output bytes=60
Map output materialized bytes=78
Input split bytes=84
Combine input records=0
Combine output records=0
Reduce input groups=3
Reduce shuffle bytes=78
Reduce input records=6
Reduce output records=1
Spilled Records=12
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=18
Total committed heap usage (bytes)=403701760

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```
Bytes Read=60
```

File Output Format Counters

```
Bytes Written=25
```

```
2025-05-24 17:20:45,936 INFO streaming.StreamJob: Output directory: /bda/out1
```

```
prajwal@PrajwalDevice:~$ hdfs dfs -cat /bda/out1/part-00000
```

```
Mean Temperature: 31.18
```

```
prajwal@PrajwalDevice:~$
```

Experiment – 7

Q) For a given Text file, Create a Map Reduce program to sort the content in an alphabetic order listing only top 10 maximum occurrences of words.

Code & Output:

Top N Words Using MapReduce

TopN.java (Driver)

java

CopyEdit

```
package samples.topn;
```

```
import java.io.IOException;
```

```
import java.util.StringTokenizer;
```

```
import org.apache.hadoop.conf.Configuration;
```

```
import org.apache.hadoop.fs.Path;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Job;
```

```
import org.apache.hadoop.mapreduce.Mapper;
```

```
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
```

```
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
import org.apache.hadoop.util.GenericOptionsParser;
```

```
public class TopN {
```

```
    public static void main(String[] args) throws Exception {
```

```
        Configuration conf = new Configuration();
```

```
        String[] otherArgs = (new GenericOptionsParser(conf, args)).getRemainingArgs();
```

```
        if (otherArgs.length != 2) {
```

```
            System.err.println("Usage: TopN <in> <out>");
```

```
System.exit(2);
}
```

```
Job job = Job.getInstance(conf);
job.setJobName("Top N");
job.setJarByClass(TopN.class);
job.setMapperClass(TopNMapper.class);
job.setReducerClass(TopNReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
```

```
FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
```

```
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

```
public static class TopNMapper extends Mapper<Object, Text, Text, IntWritable> {
private static final IntWritable one = new IntWritable(1);
private Text word = new Text();
private String tokens = "[_#$%&^`=\\[\\]\\\\\\*\\/\\\\\\\\,;\\.\\-:()?!\"'"]";

public void map(Object key, Text value, Mapper<Object, Text, Text, IntWritable>.Context context)
throws IOException, InterruptedException {
String cleanLine = value.toString().toLowerCase().replaceAll(this.tokens, " ");
StringTokenizer itr = new StringTokenizer(cleanLine);

while (itr.hasMoreTokens()) {
this.word.set(itr.nextToken().trim());
context.write(this.word, one);
}
```

```

}
}
}
}

```

TopNCombiner.java

```
java
```

```
CopyEdit
```

```
package samples.topn;
```

```
import java.io.IOException;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Reducer;
```

```
public class TopNCombiner extends Reducer<Text, IntWritable, Text, IntWritable> {
```

```
public void reduce(Text key, Iterable<IntWritable> values,
```

```
Reducer<Text, IntWritable, Text, IntWritable>.Context context)
```

```
throws IOException, InterruptedException {
```

```
int sum = 0;
```

```
for (IntWritable val : values)
```

```
sum += val.get();
```

```
context.write(key, new IntWritable(sum));
```

```
}
```

```
}
```

TopNMapper.java

```
java
```

```
CopyEdit
```

```
package samples.topn;
```

```

import java.io.IOException;

import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class TopNMapper extends Mapper<Object, Text, Text, IntWritable> {
    private static final IntWritable one = new IntWritable(1);
    private Text word = new Text();
    private String tokens = "[_!$#<>\\^=\\[\\]\\|\\*\\/\\\\\\,;,.\\|-:()?!\"'"]";

    public void map(Object key, Text value, Mapper<Object, Text, Text, IntWritable>.Context context)
    throws IOException, InterruptedException {
        String cleanLine = value.toString().toLowerCase().replaceAll(this.tokens, " ");
        StringTokenizer itr = new StringTokenizer(cleanLine);

        while (itr.hasMoreTokens()) {
            this.word.set(itr.nextToken().trim());
            context.write(this.word, one);
        }
    }
}

```

TopNReducer.java

```

java
CopyEdit
package samples.topn;

import java.io.IOException;

```

```

import java.util.HashMap;
import java.util.Map;
import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
import utils.MiscUtils;

public class TopNReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    private Map<Text, IntWritable> countMap = new HashMap<>();

    public void reduce(Text key, Iterable<IntWritable> values,
        Reducer<Text, IntWritable, Text, IntWritable>.Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values)
            sum += val.get();
        this.countMap.put(new Text(key), new IntWritable(sum));
    }

    protected void cleanup(Reducer<Text, IntWritable, Text, IntWritable>.Context context)
        throws IOException, InterruptedException {
        Map<Text, IntWritable> sortedMap = MiscUtils.sortByValues(this.countMap);
        int counter = 0;
        for (Text key : sortedMap.keySet()) {
            if (counter++ == 20)
                break;
            context.write(key, sortedMap.get(key));
        }
    }
}

```

Experiment – 8

Q) Write a Scala program to print numbers from 1 to 100 using for loop.

Code

```
object PrintNumbers {  
  def main(args: Array[String]): Unit = {  
    for (i <- 1 to 100) {  
      println(i)  
    }  
  }  
}
```

Output:

```
scala> for(i <- 0 to 100){  
  | println(i)  
  | }  
0  
1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34
```

Experiment – 9

Q) Using RDD and FlatMap count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark

Code:

```
from pyspark.sql import SparkSession

from pyspark.sql.functions import udf, col, trim, lower

from pyspark.sql.types import ArrayType, StringType

import nltk

from nltk.corpus import stopwords

from nltk.stem import WordNetLemmatizer


# Download nltk data (run once)

nltk.download('stopwords')

nltk.download('wordnet')


# Initialize Spark session

spark = SparkSession.builder \

    .appName("SimpleTextStreamingCleaning") \

    .getOrCreate()


# Set log level to WARN to reduce verbosity

spark.sparkContext.setLogLevel("WARN")


# Define host and port to listen for streaming text data

host = "localhost"

port = 9999

# Read streaming data from socket

lines = spark.readStream.format("socket") \
```



```

.option("host", host) \
.option("port", port) \
.load()
# Convert each line to lowercase and trim whitespace
lines_cleaned = lines.select(trim(lower(col("value"))).alias("line"))
# Define stop words set
stop_words_set = set(stopwords.words('english'))
# Initialize lemmatizer
lemmatizer = WordNetLemmatizer()
# Define UDF for tokenization, stop words removal, and lemmatization
def clean_text(line):
    if not line:
        return []
    tokens = line.split()
    tokens = [word for word in tokens if word not in stop_words_set]
    lemmas = [lemmatizer.lemmatize(word) for word in tokens]
    return lemmas
clean_text_udf = udf(clean_text, ArrayType(StringType()))
# Apply cleaning UDF to each line
cleaned = lines_cleaned.withColumn("cleaned_tokens", clean_text_udf(col("line")))
# Convert tokens back to string for display
from pyspark.sql.functions import concat_ws
final_output = cleaned.select(concat_ws(" ", col("cleaned_tokens")).alias("cleaned_line"))

# Write stream to console
query = final_output.writeStream \
    .outputMode("append") \
    .format("console") \

```

```
.option("truncate", False) \
```

```
.start()
```

```
query.awaitTermination()
```

Output:

```
prajwal@ppc-jna1@work:~$ spark-shell
25/05/24 17:41:38 WARN Utils: Your hostname, PrajwalDevice resolves to a loopback address: 127.0.1.1; using 10.255.255.254 instead (on interface lo)
25/05/24 17:41:38 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
25/05/24 17:41:46 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Spark context Web UI available at http://10.255.255.254:4040
Spark context available as 'sc' (master = local[*], app id = local-1748888707553).
Spark session available as 'spark'.
Welcome to

  ____              __
 / ___|  _ \  _ __ | | | |
 \___ \ (_) | | '_ \| | | |
  ___) / ___/| | |_) | | | |
 |____/_|___||_|.___||_|_|_|

version 3.5.5

Using Scala version 2.12.18 (OpenJDK 64-Bit Server VM, Java 21.0.7)
Type in expressions to have them evaluated.
Type :help for more information.

scala> val file=sc.text25/05/24 17:42:00 WARN GarbageCollectionMetrics: To enable non-built-in garbage collector(s) List(G1 Concurrent GC), users should configure it(them) to spark.e
spark.eventlog.gcMetrics.oldGenerationGarbageCollectors
val file=sc.textFile("il.txt")
file: org.apache.spark.rdd.RDD[String] = il.txt MapPartitionsRDD[1] at textFile at <console>:23

scala> val words=file.flatMap(line=>line.split("\\W+"))
words: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[2] at flatMap at <console>:23

scala> val wordpairs=words.map(word=>{word.toLowerCase,1})
wordpairs: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[3] at map at <console>:23

scala> val wordc=wordpairs.reduceByKey(_+_.)
wordc: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceByKey at <console>:23

scala> val fil2=wordc.filter{case(word,count)>=>count>2}
fil2: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[5] at filter at <console>:23

scala> fil2.collect().foreach(println)
```